GPGPU Programming

Why GP-GPU?

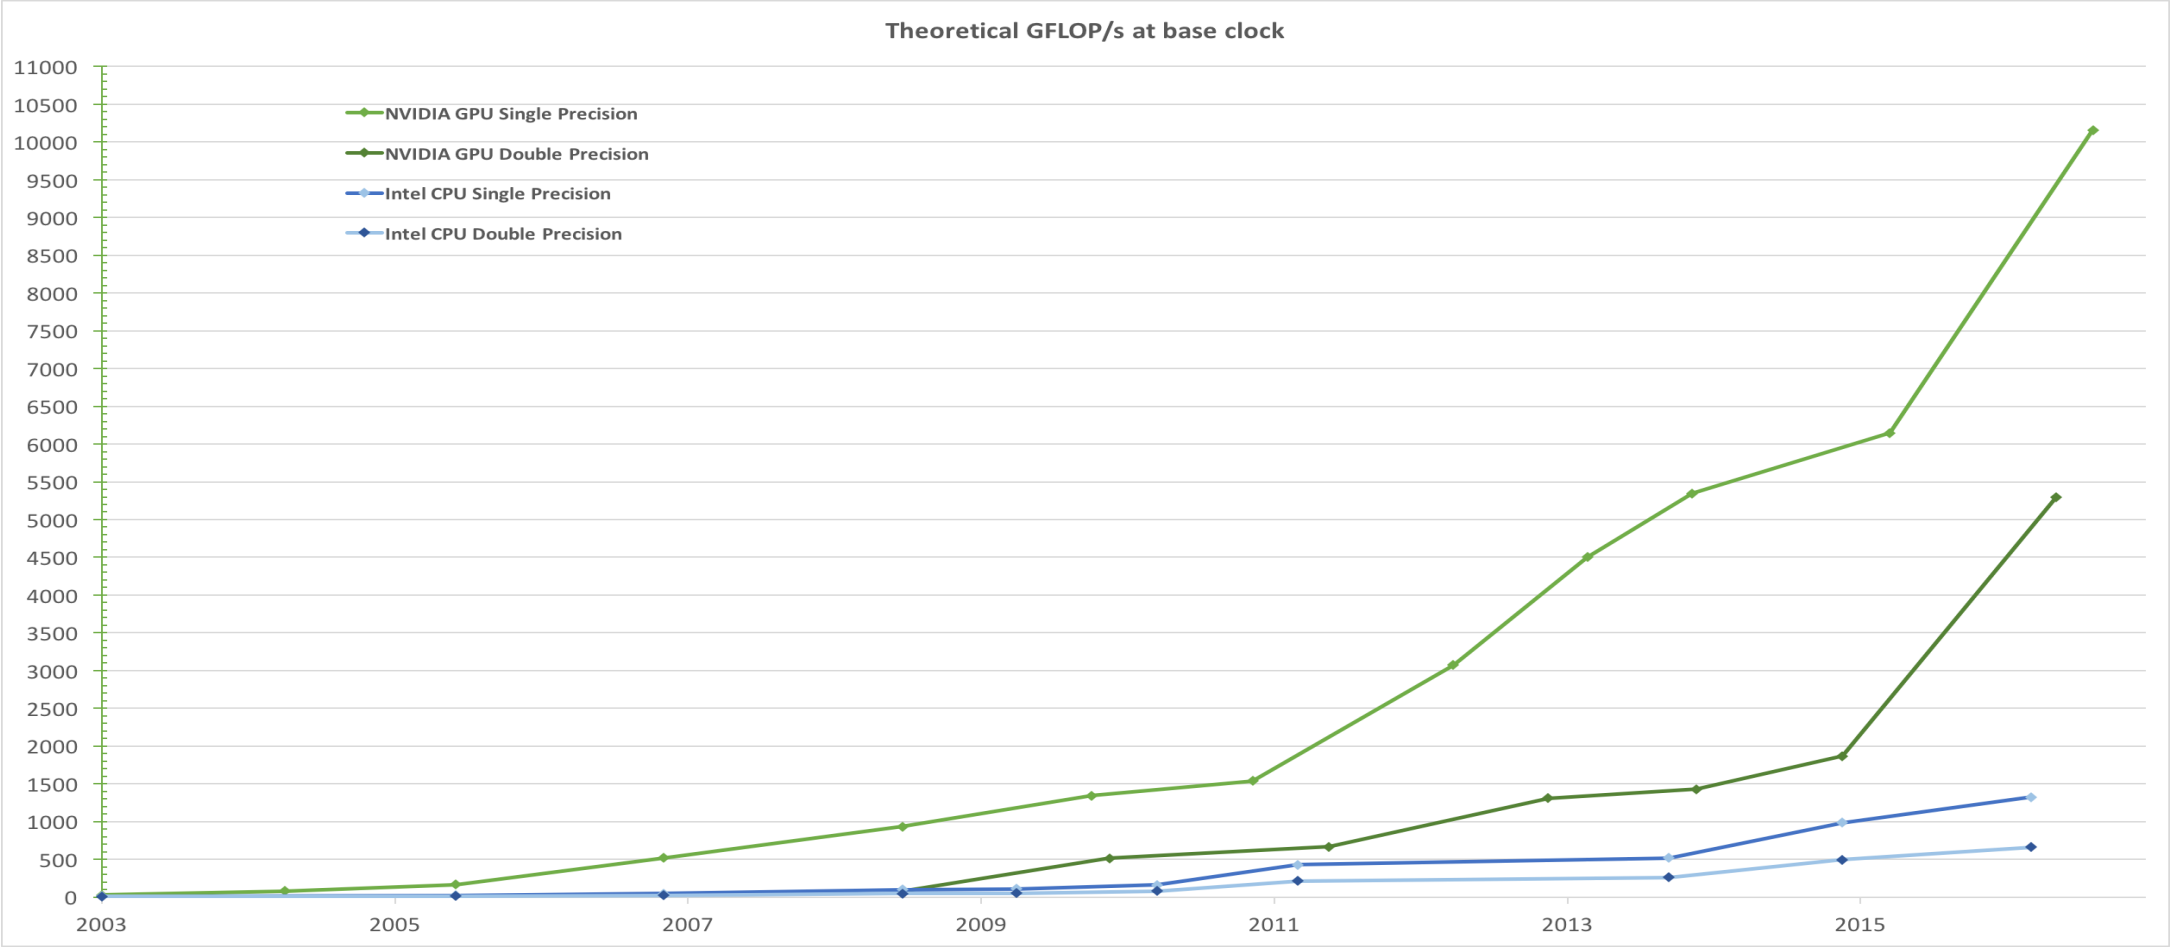

Source: CUDA C Programming Guide (NVIDIA)

Why GP-GPU?

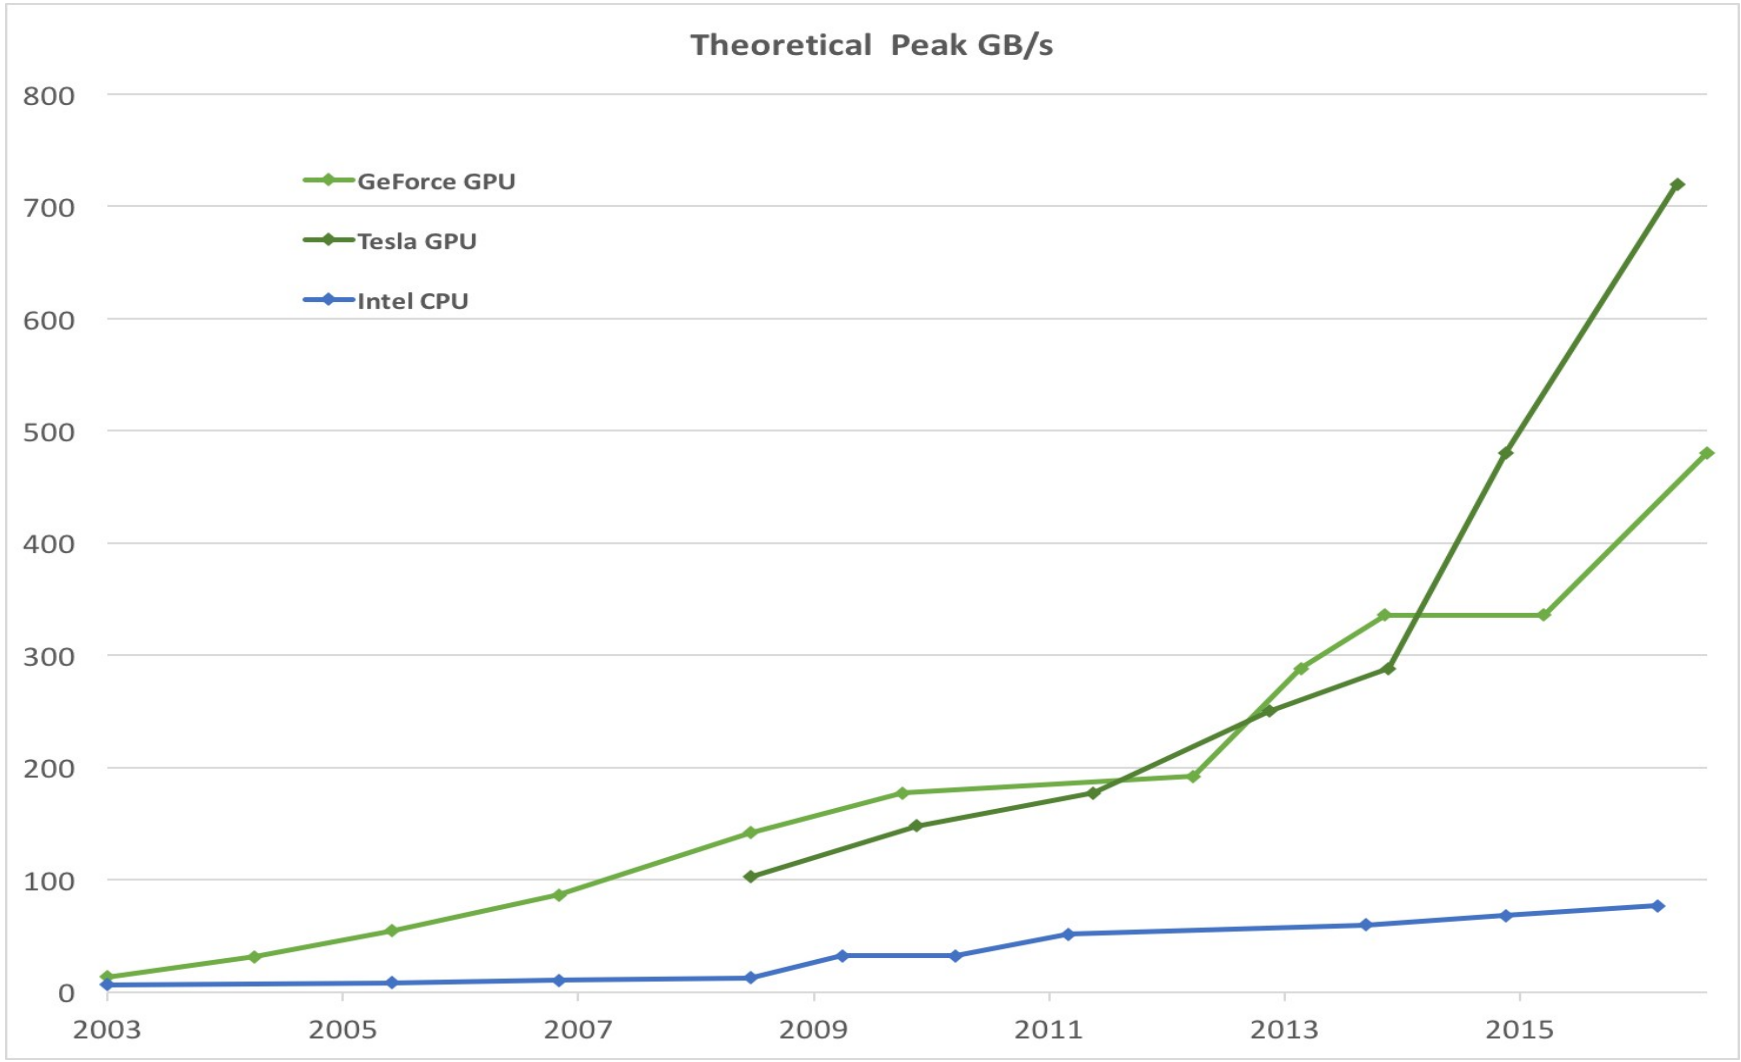

Source: CUDA C Programming Guide (NVIDIA)

Where is my GPU?

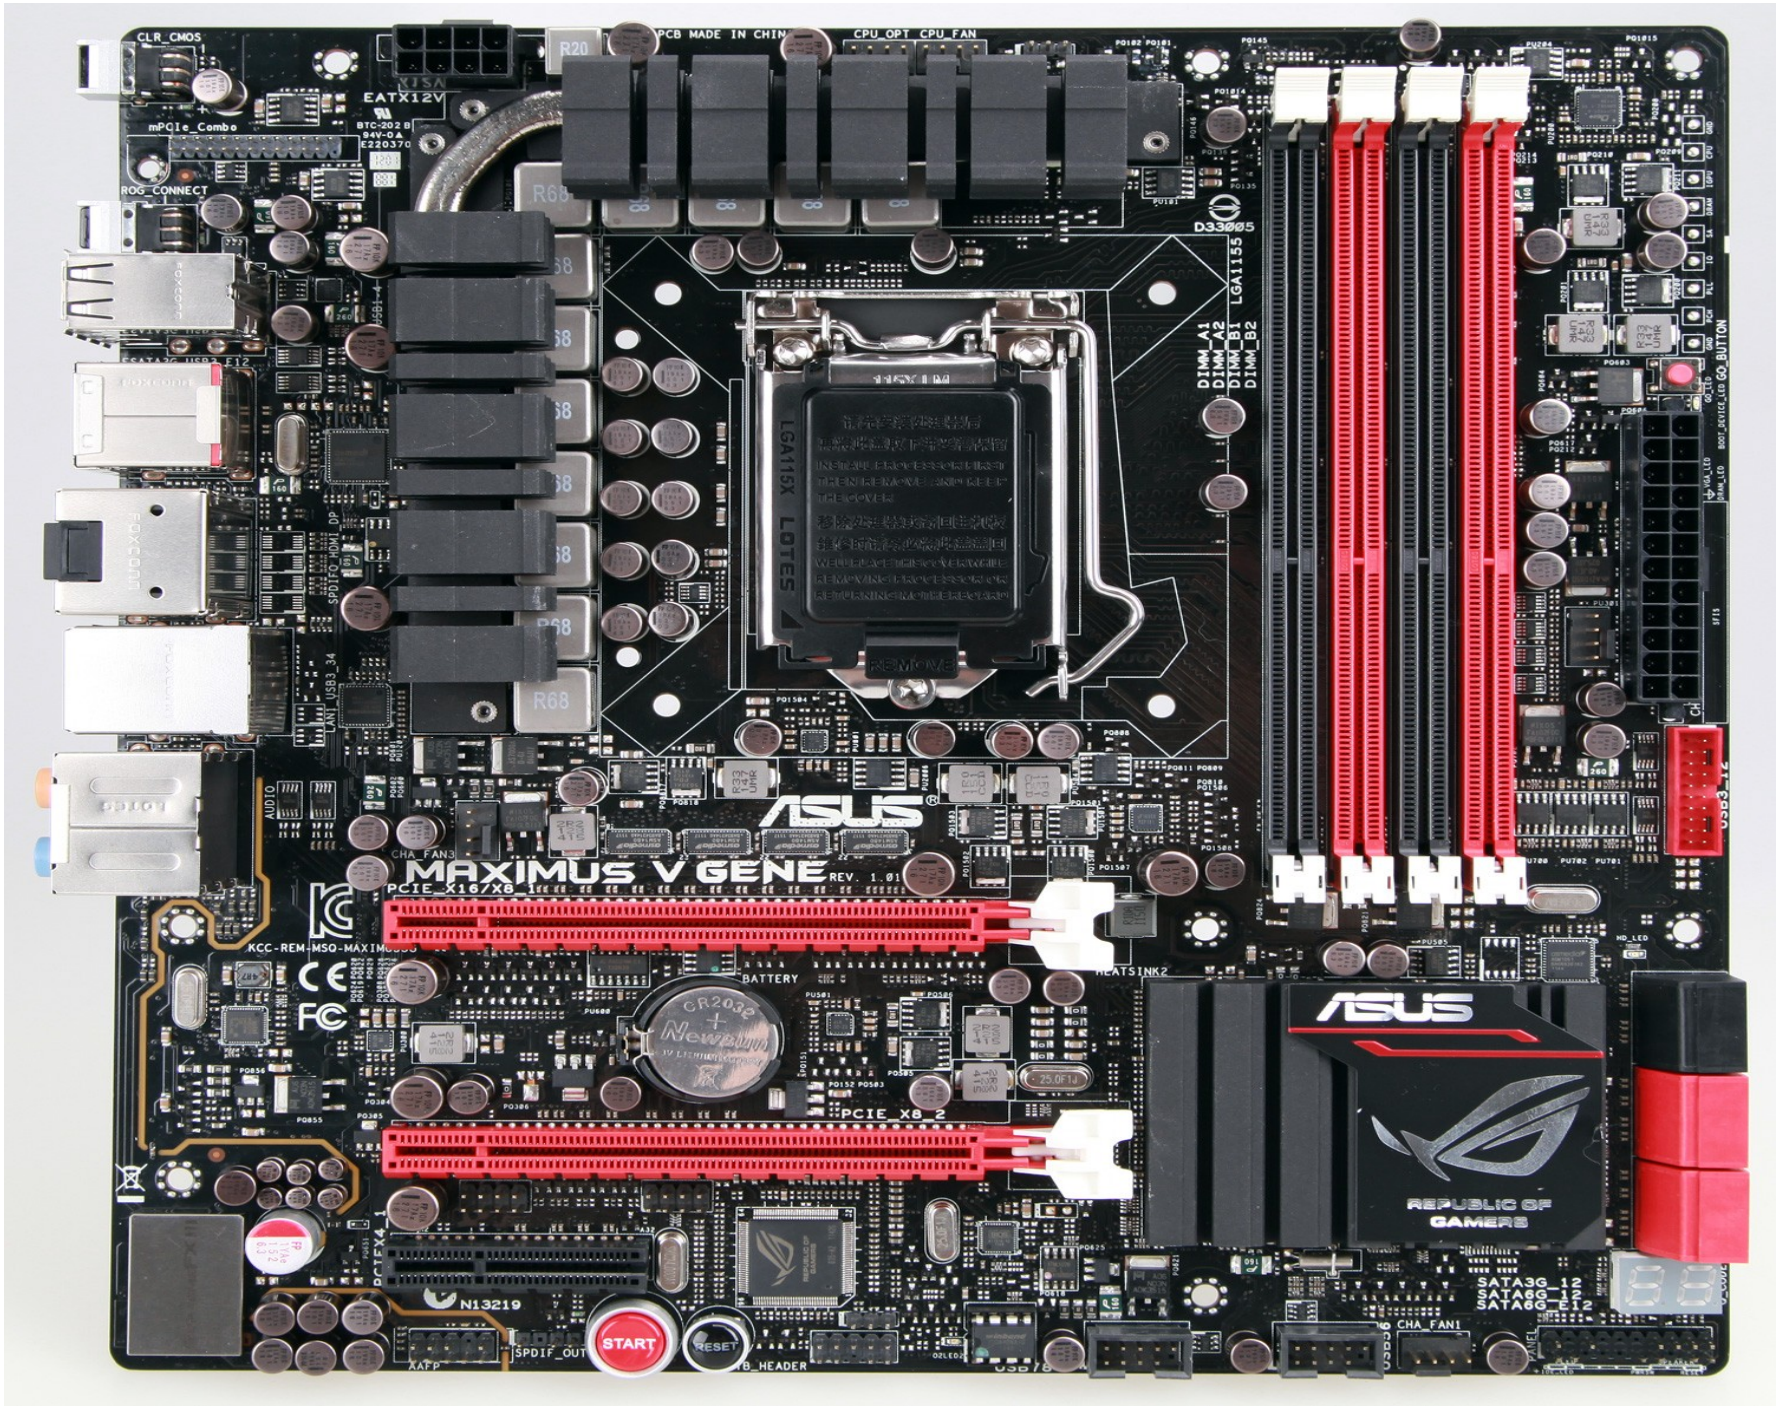
CLR_CMOS
EATX12V
mPCIe_Combo
ROG_CONNECT
PCB MADE IN CHINA
CPU_OPT
CPU_FAN
D33005
DIMM_A1
DIMM_A2
DIMM_B1
DIMM_B2
LGA1155
LGA115X
LOTES
REMOVE
ASUS
MAXIMUS V GENE
REV. 1.01
PCIE_X16/X8_1
PCIE_X8_2
BATTERY
CR2032
Newsun
HEATSINK2
REPUBLIC OF GAMERS
SATA3G_12
SATA6G_12
SATA6G_E12
CHA_FAN1
CHA_FAN3
AUDIO
SPDIF_OUT
START
RESET
N13219

www.asus.com

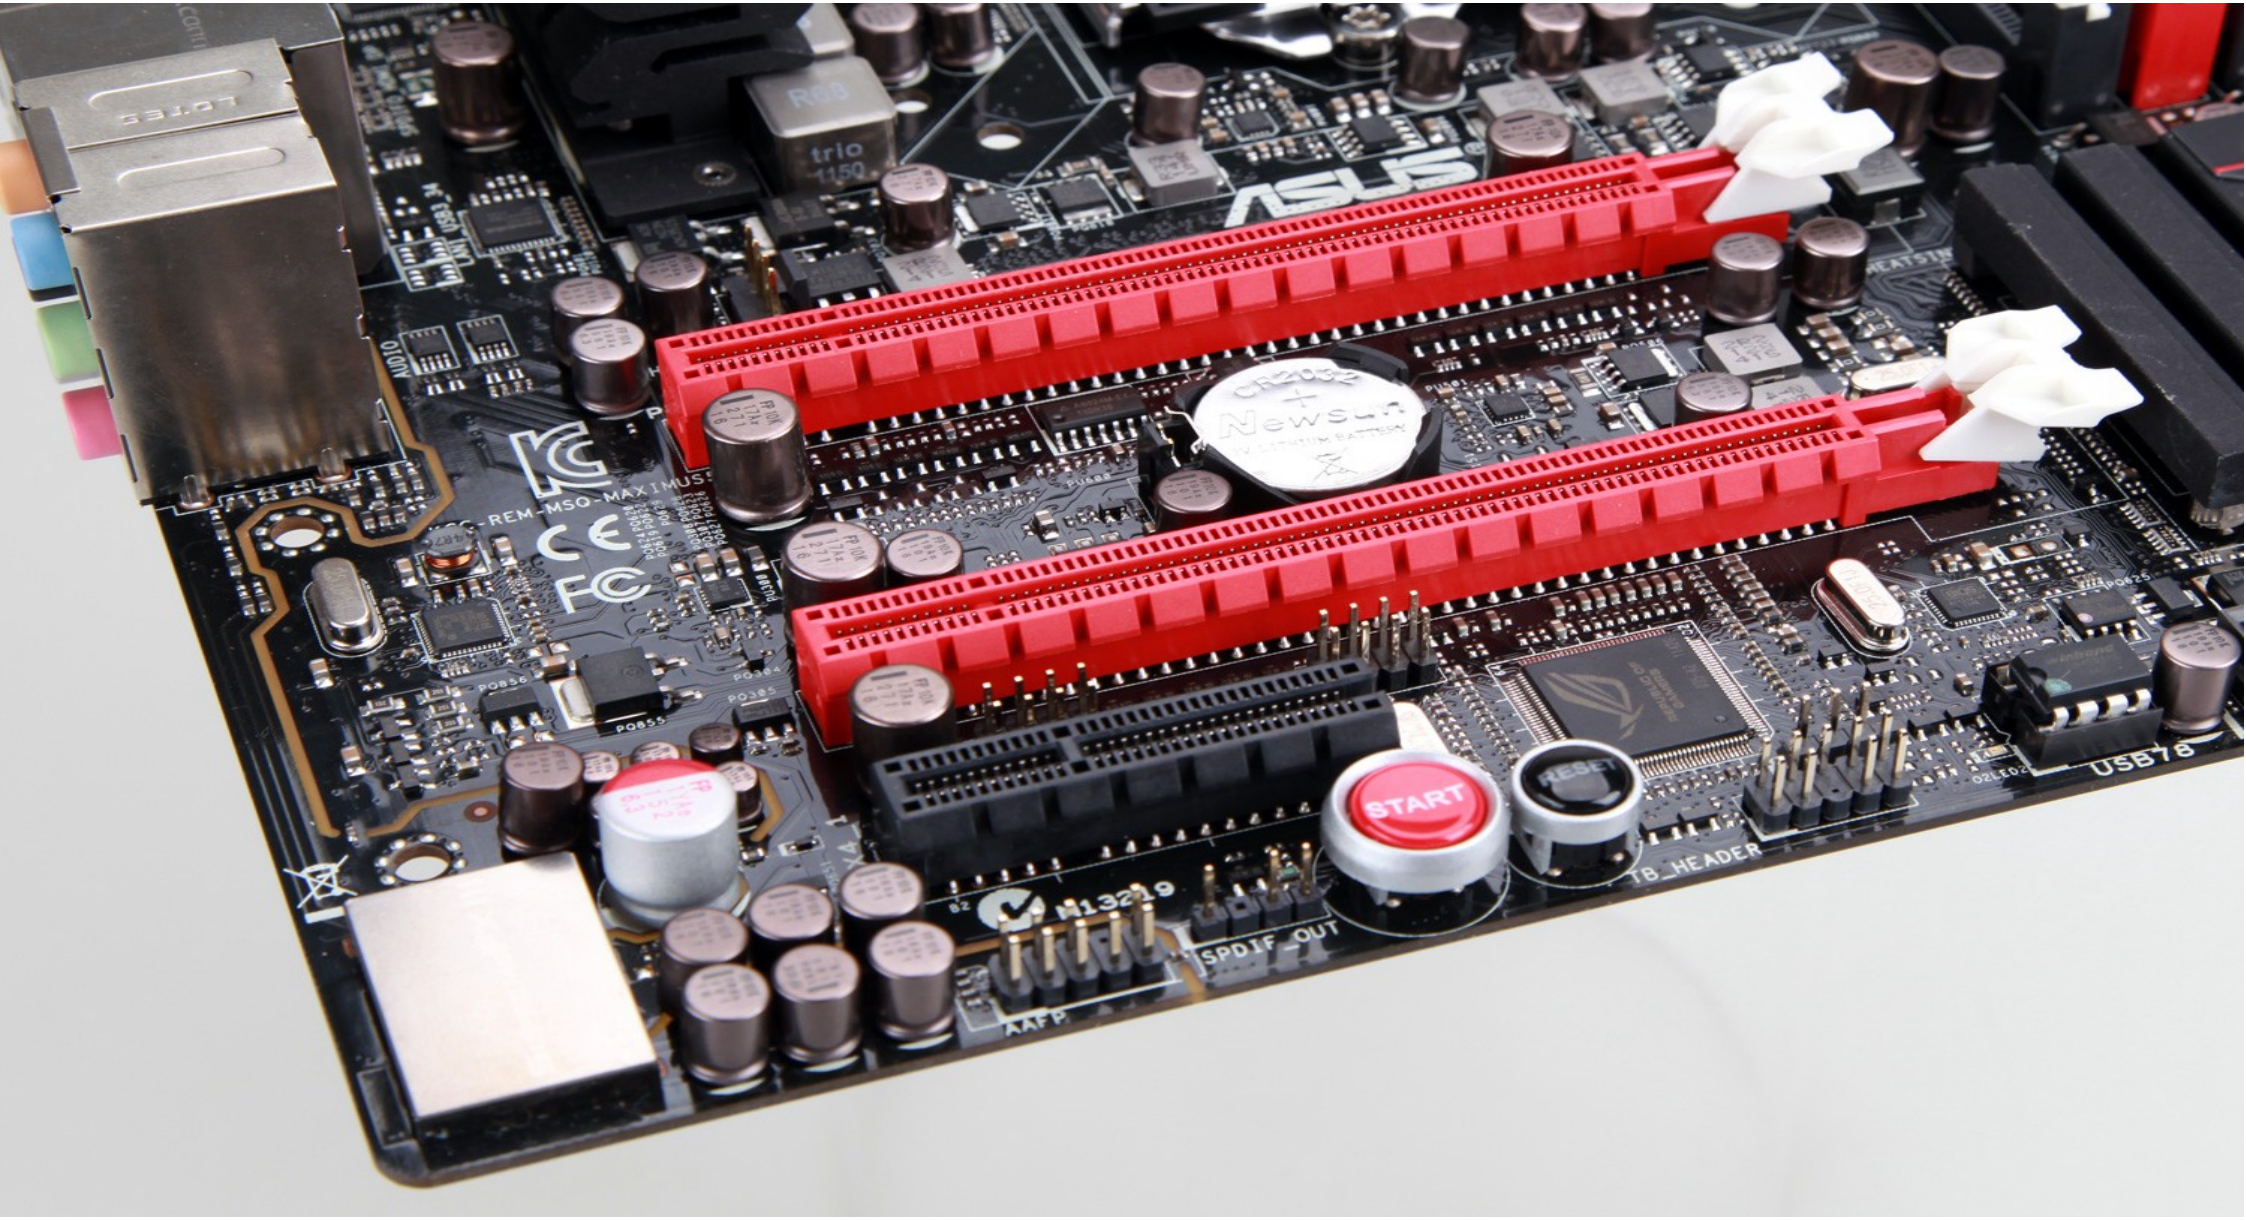

www.asus.com

www.asus.com

www.asus.com

How it differs from CPU?

CPU vs GPU

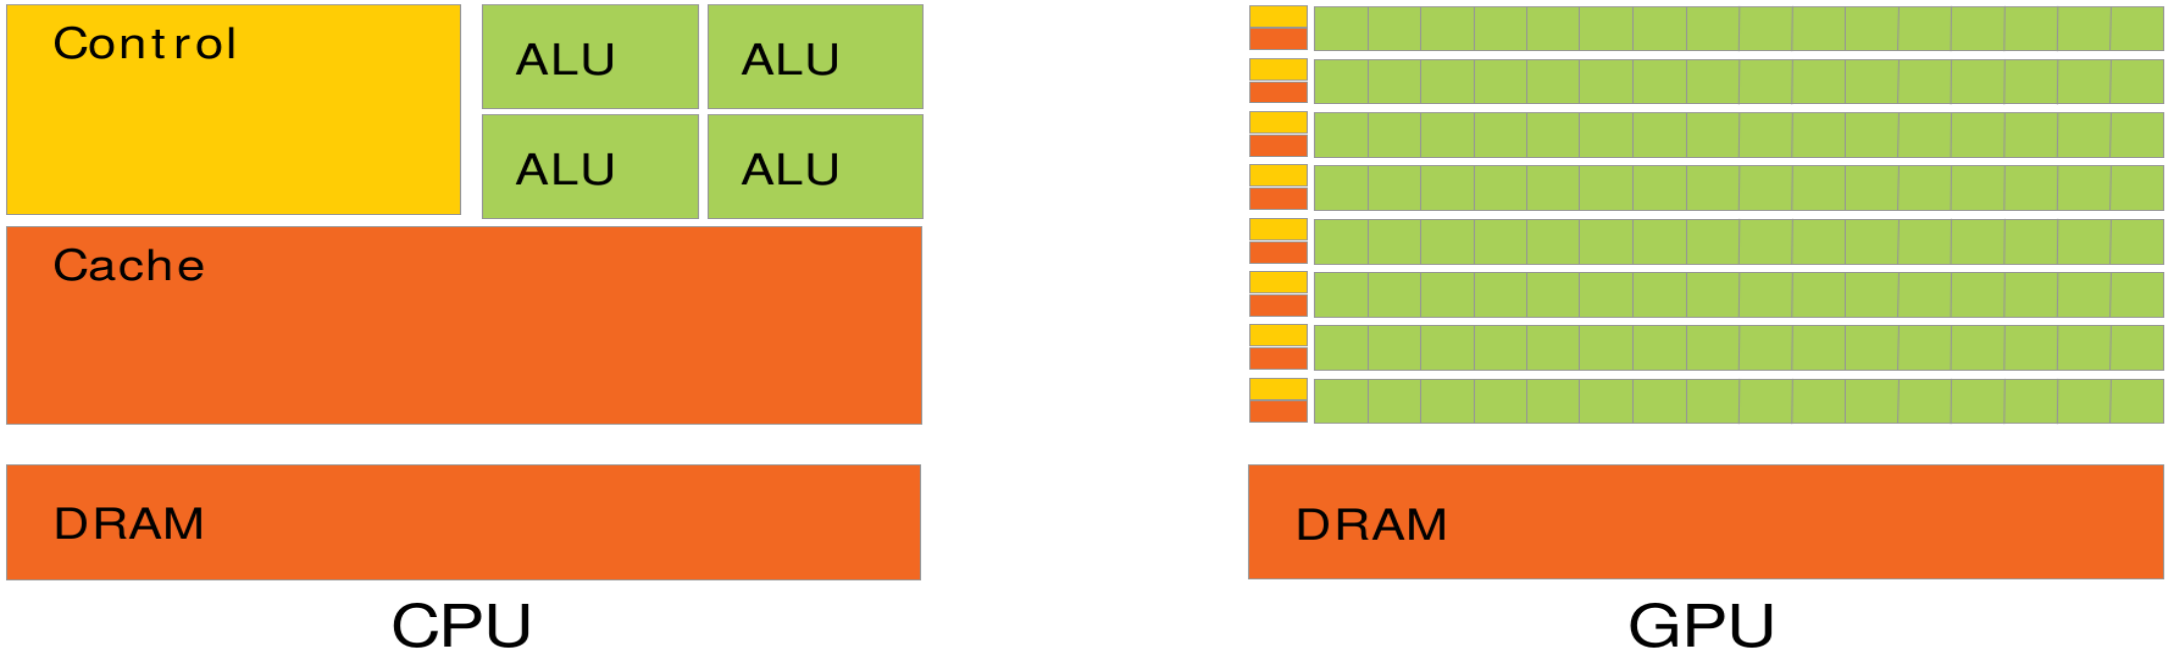

Source: CUDA C Programming Guide (NVIDIA)

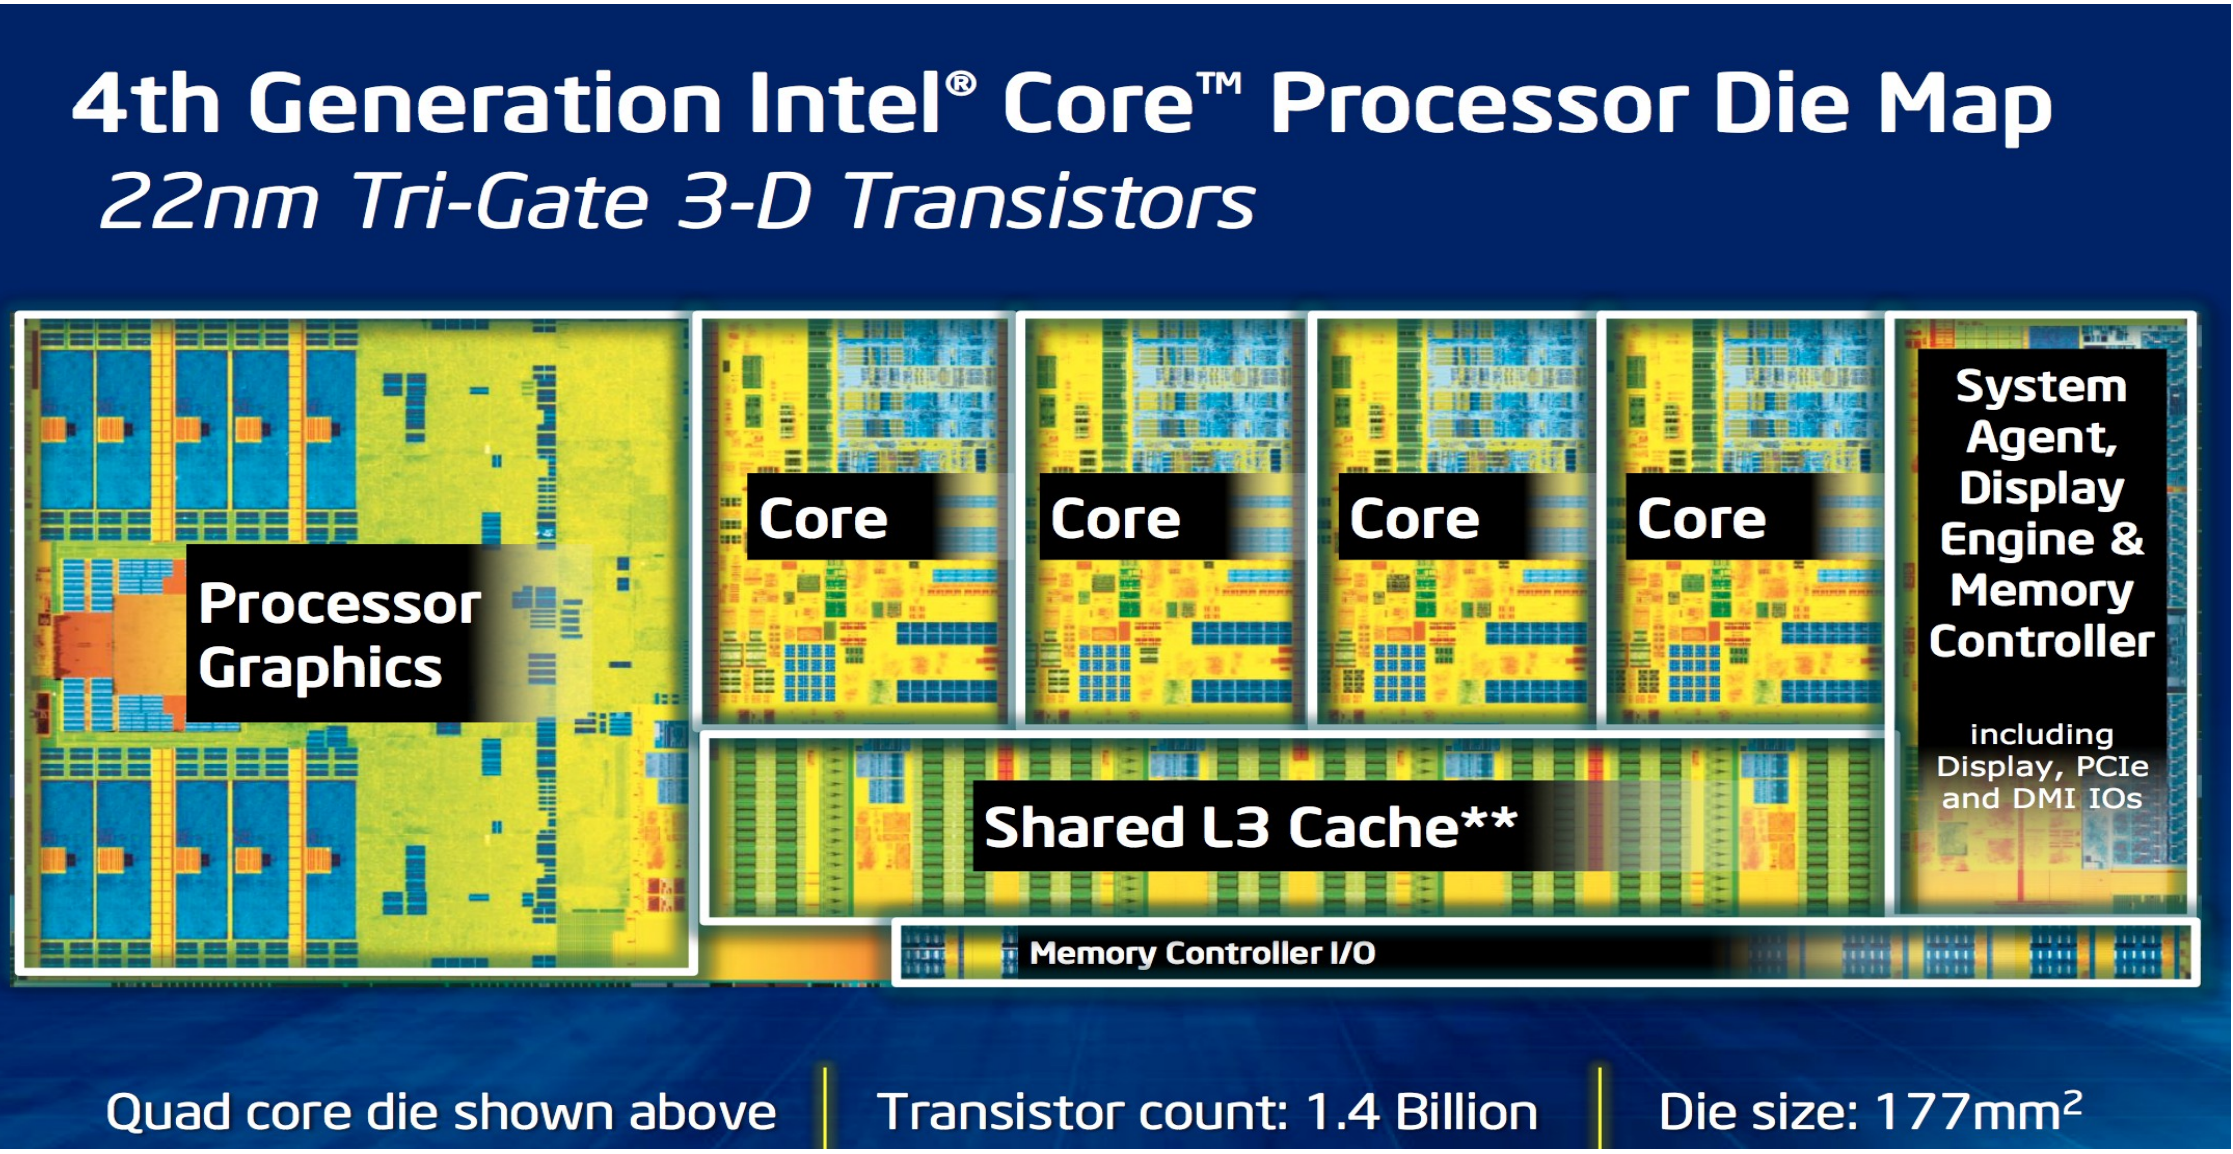

http://images.anandtech.com/doci/7003/Screen%20Shot%202013-05-31%20at%207.59.16%20PM.png

CPU vs GPU

- CPU
  - Optimized for serial thread performance
  - Good for complex tasks
  - Few, large, complex cores
  - Large number of transistors are allocated for Caches, Instruction Level Parallelism
- GPU
  - Optimized for data parallel, throughput computations
  - Large number of very small, simple cores
  - More number of transistors are allocated for computations

How to program these GPUs for general purpose computing?

Programmer’s View

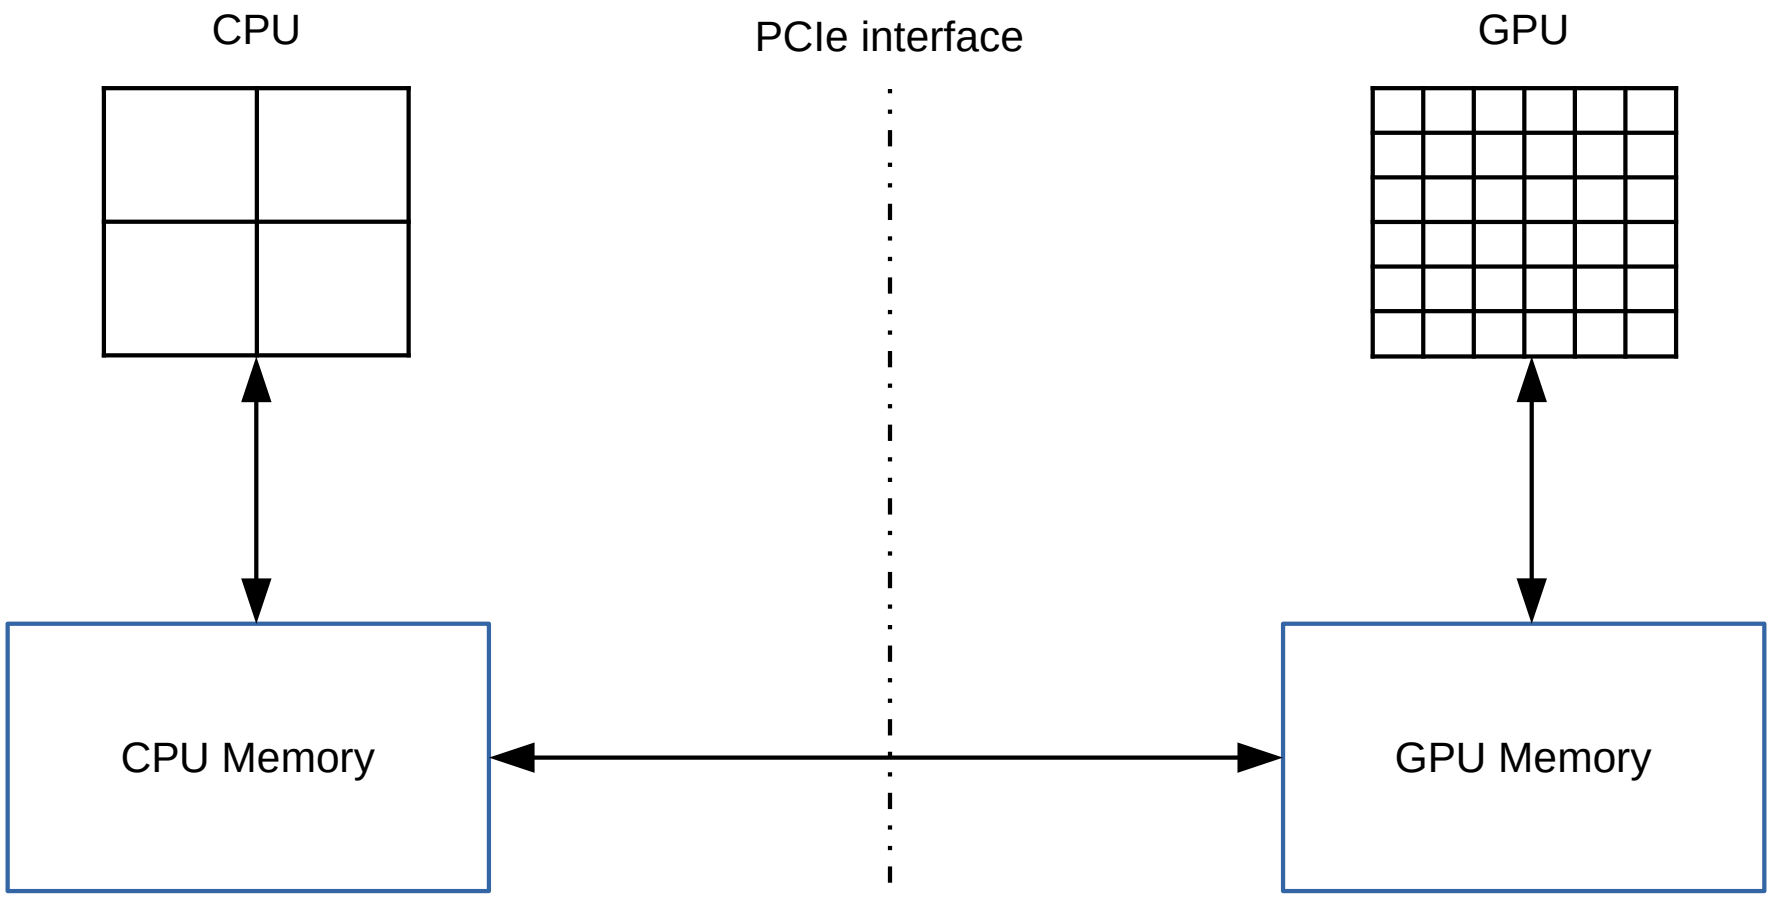

Programming GPUs

- CUDA
- OpenCL
- OpenMP 4.0+
- OpenACC
- . . .

Pseudo Code (old method)

1. Allocate and initialize memory on Host (CPU)
2. Allocate memory on Device (GPU)
3. Transfer data from Host memory to Device memory
4. Launch “kernel” on the Device – large number of threads execute in parallel on Device
5. Transfer results from Device memory to Host memory
6. De-allocate all the memory and terminate the program

Step 1

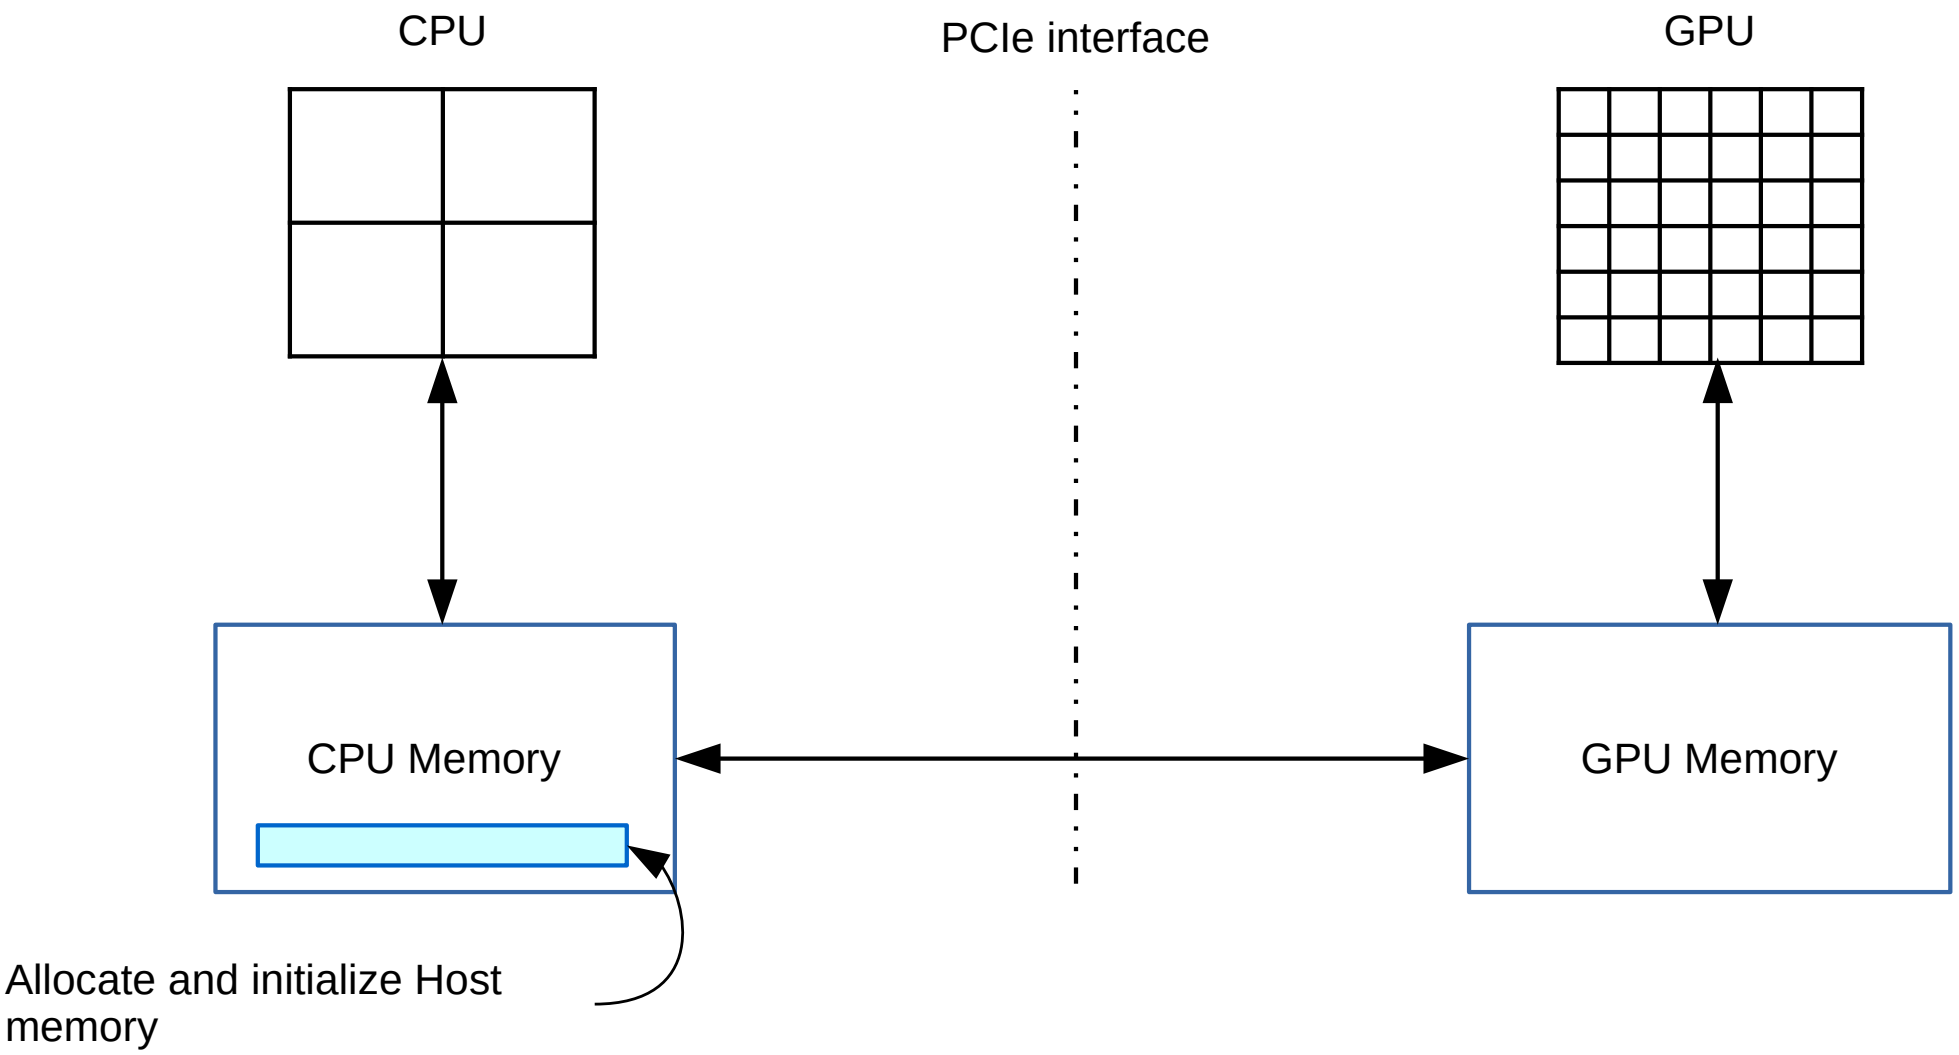
CPU
PCIe interface
GPU
CPU Memory
GPU Memory
Allocate and initialize Host memory

Step 2

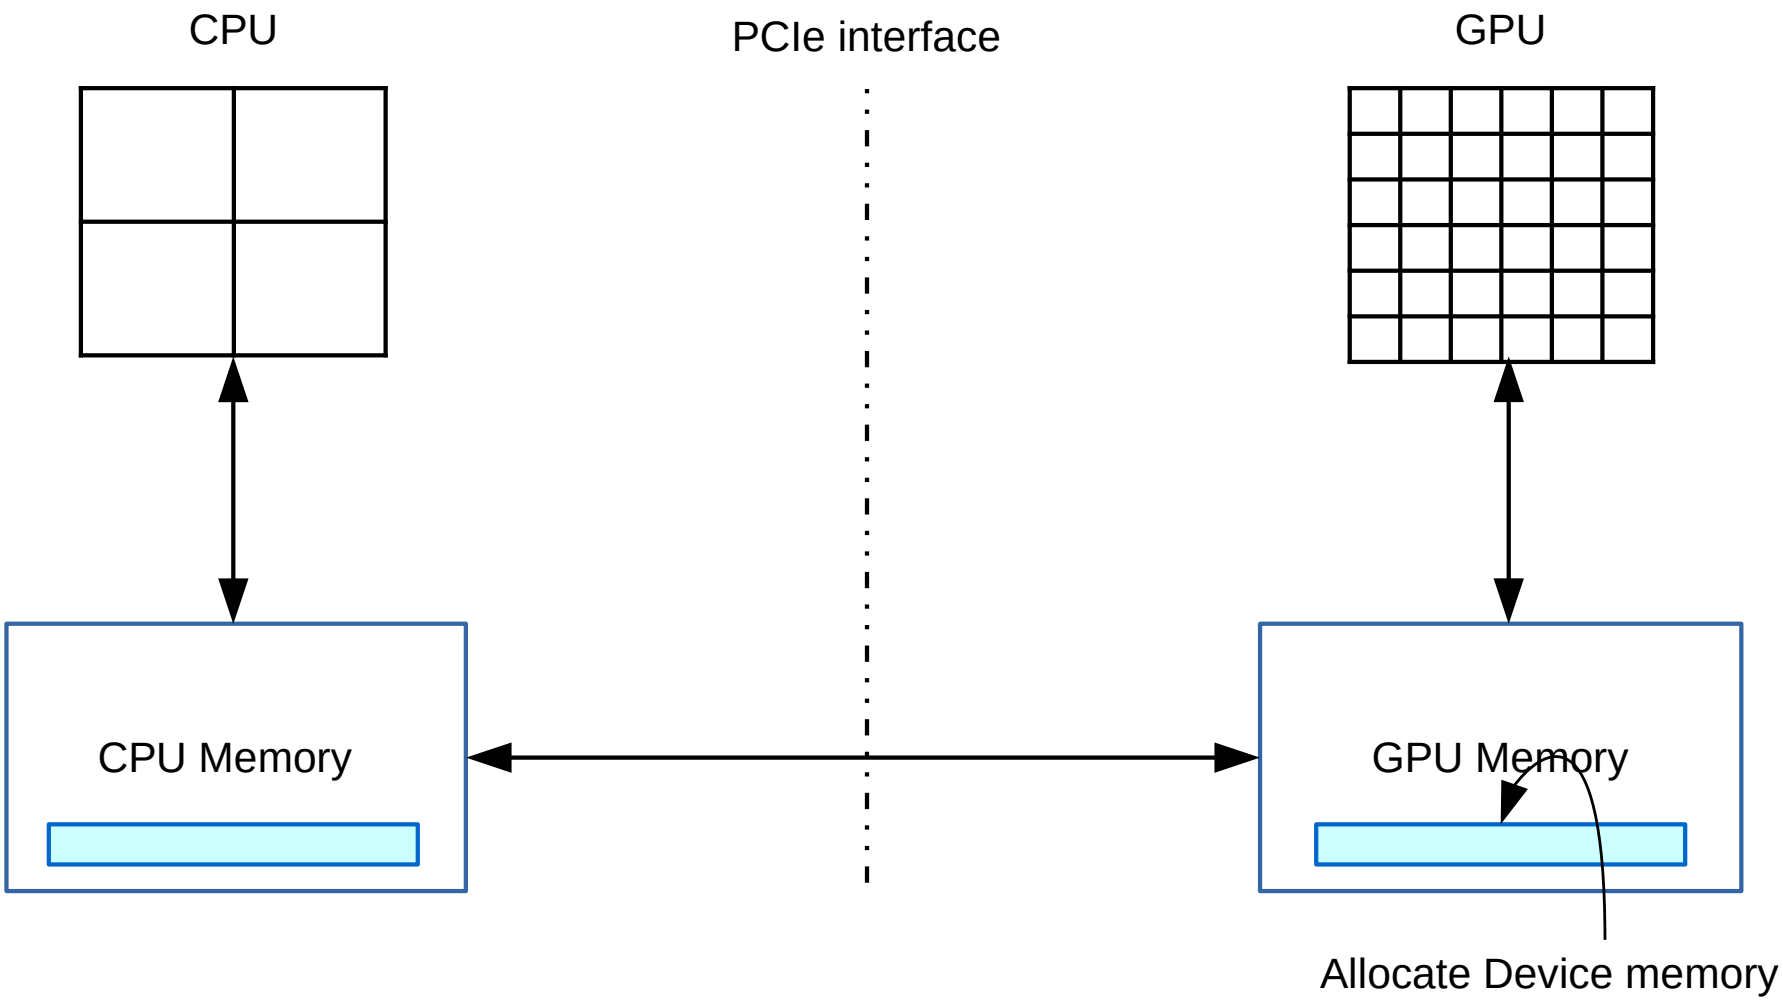
CPU
PCIe interface
GPU
CPU Memory
GPU Memory
Allocate Device memory

Step 3

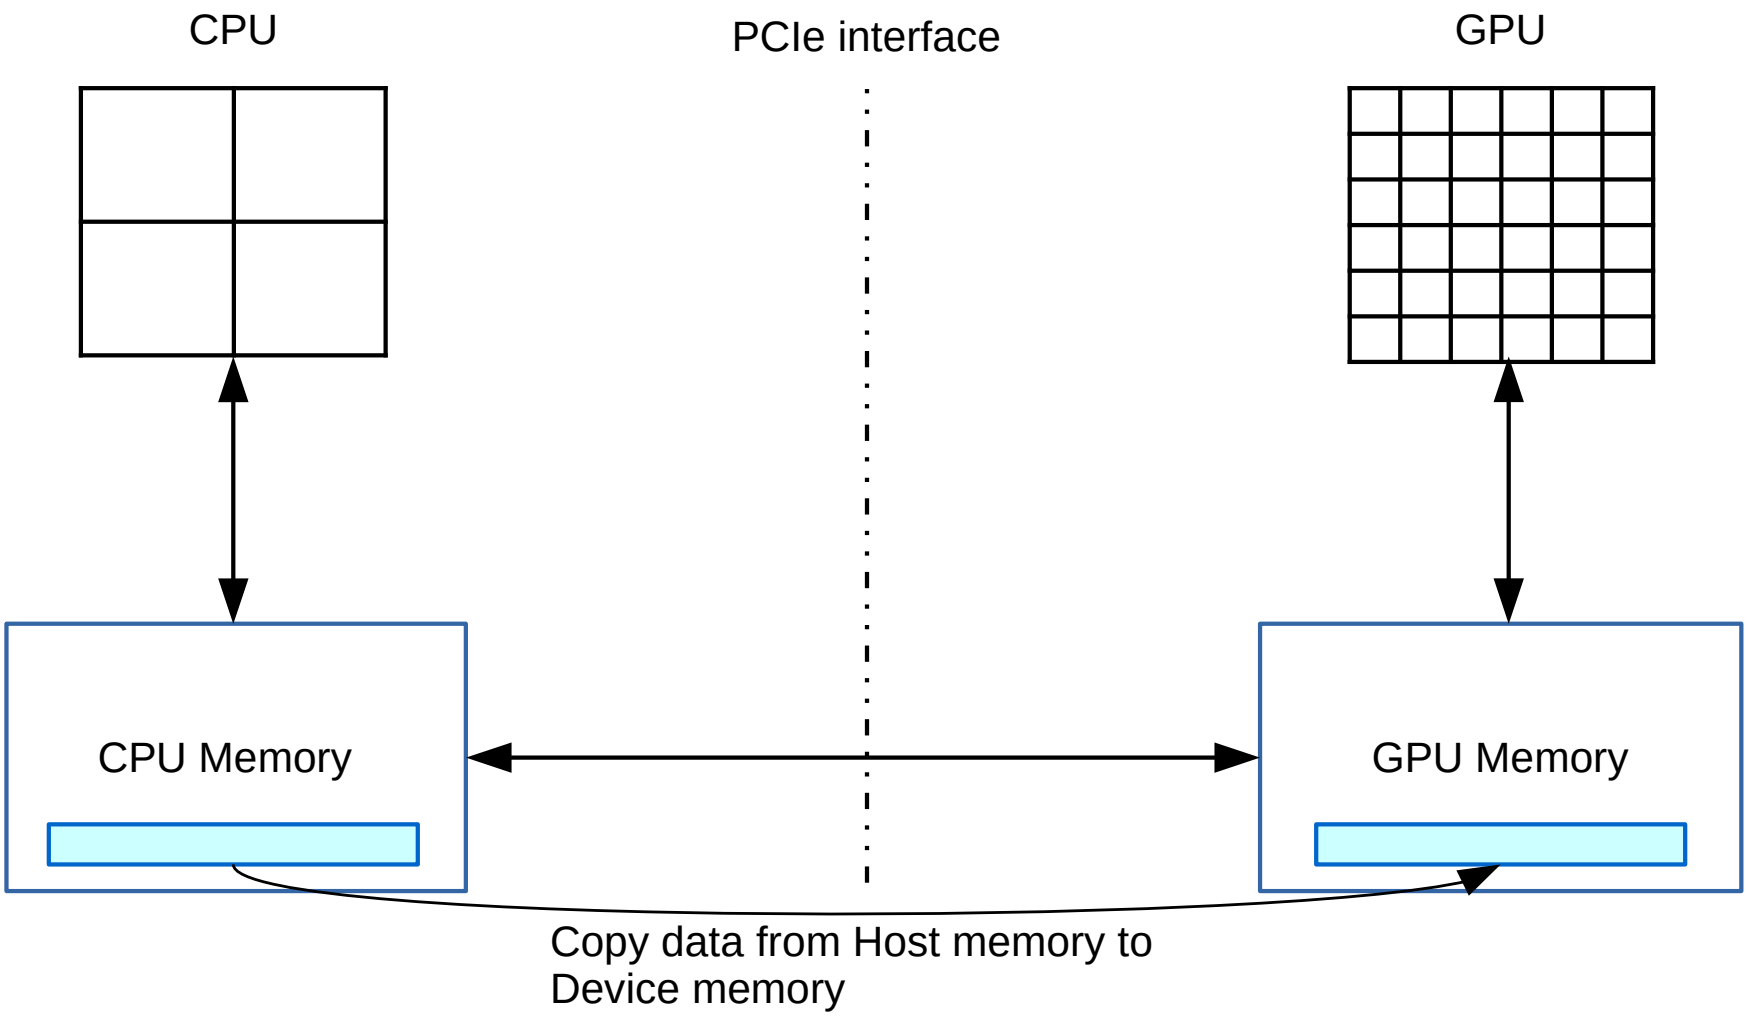
CPU
PCIe interface
GPU
CPU Memory
GPU Memory
Copy data from Host memory to
Device memory

Step 4

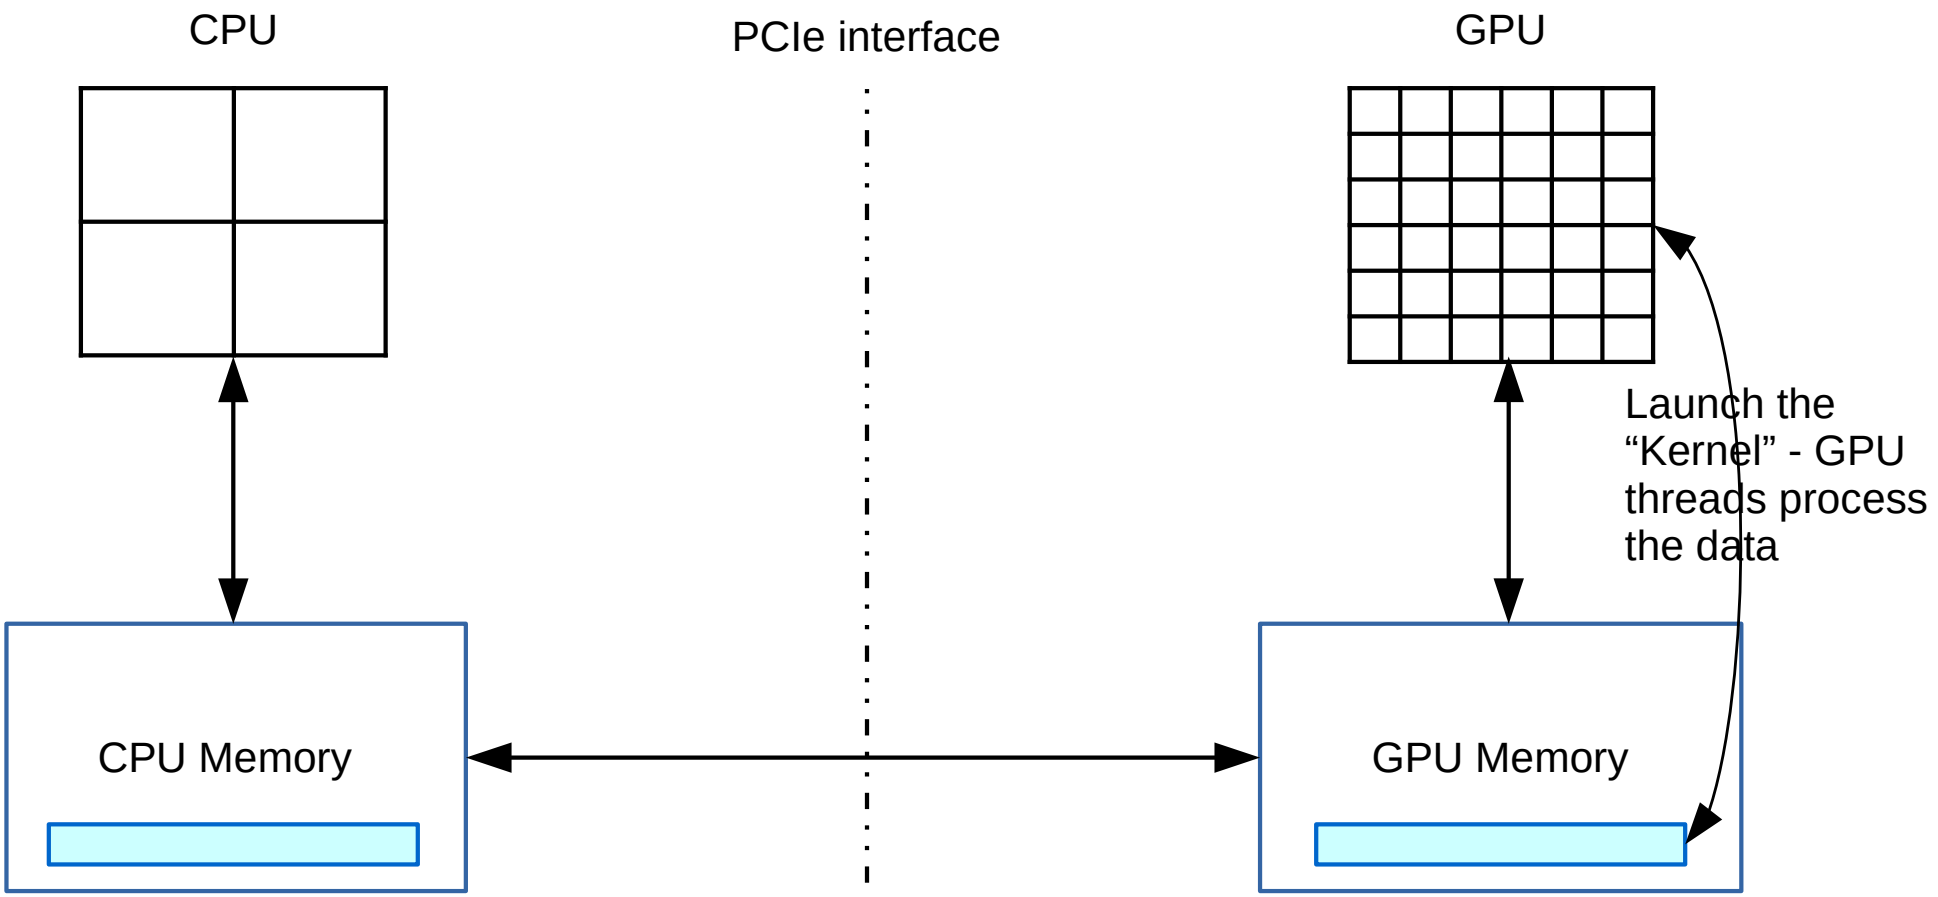
CPU
PCIe interface
GPU
Launch the
“Kernel” - GPU
threads process
the data
CPU Memory
GPU Memory

Step 5

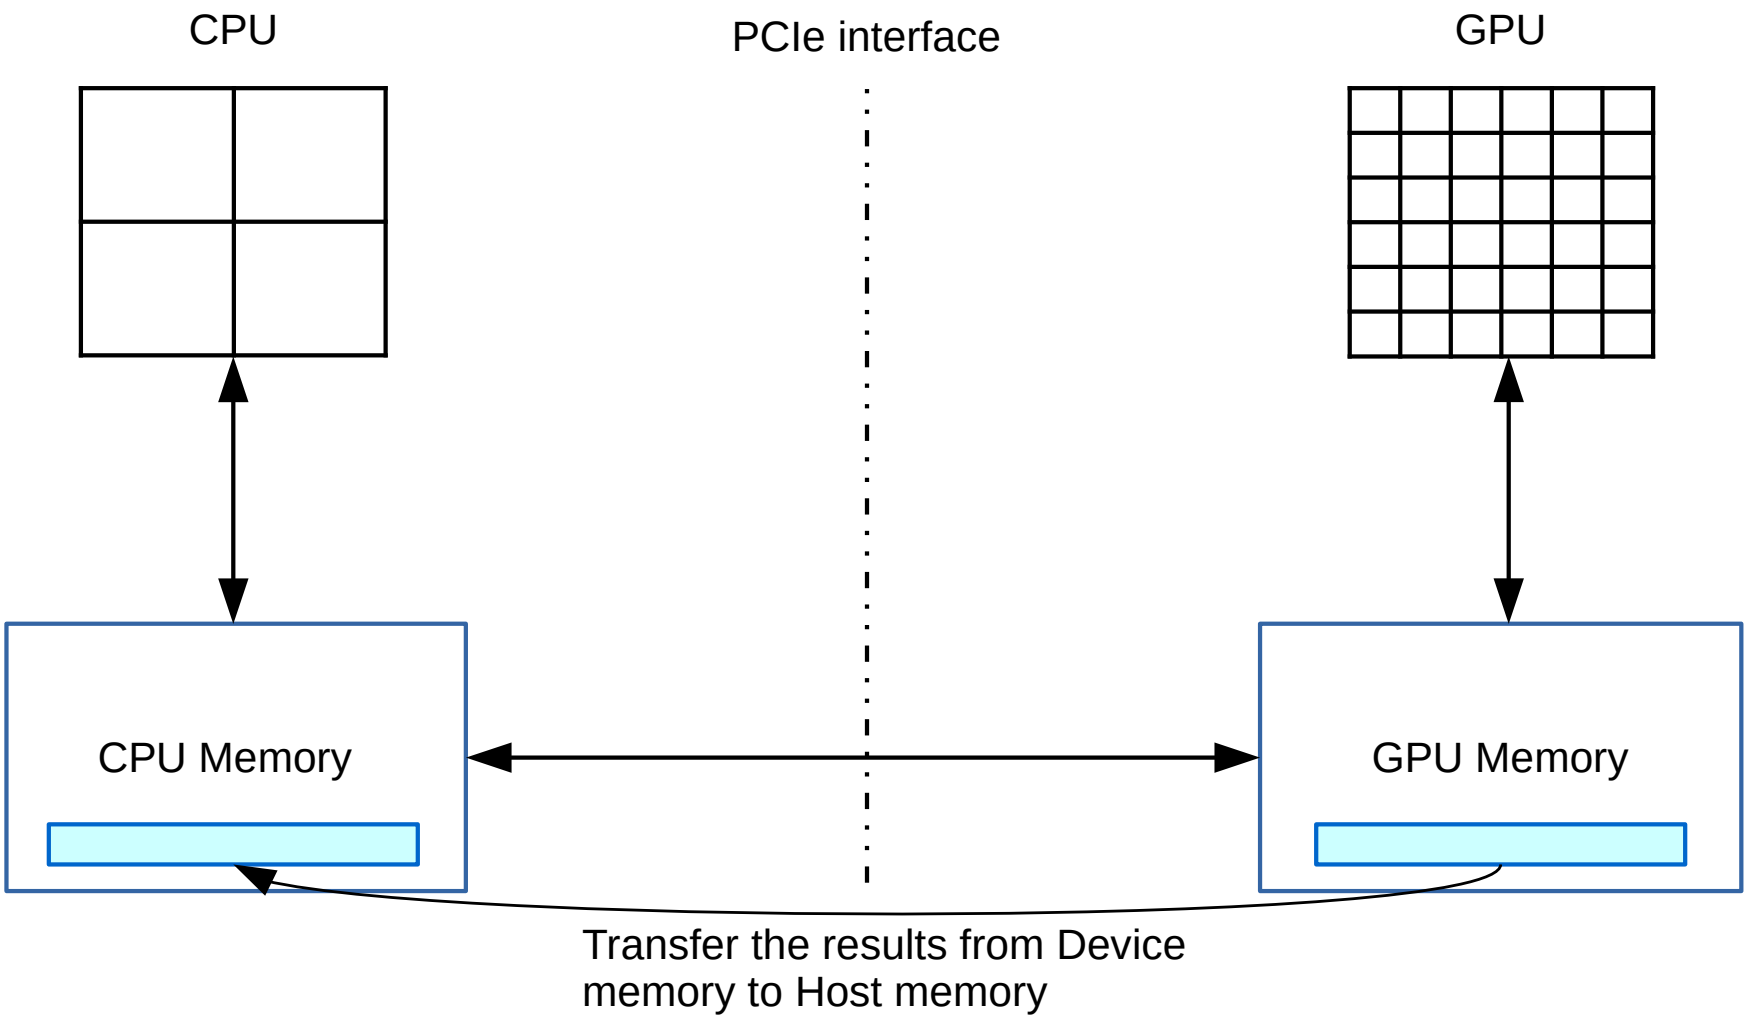
CPU
PCIe interface
GPU
CPU Memory
GPU Memory
Transfer the results from Device memory to Host memory

Step 6

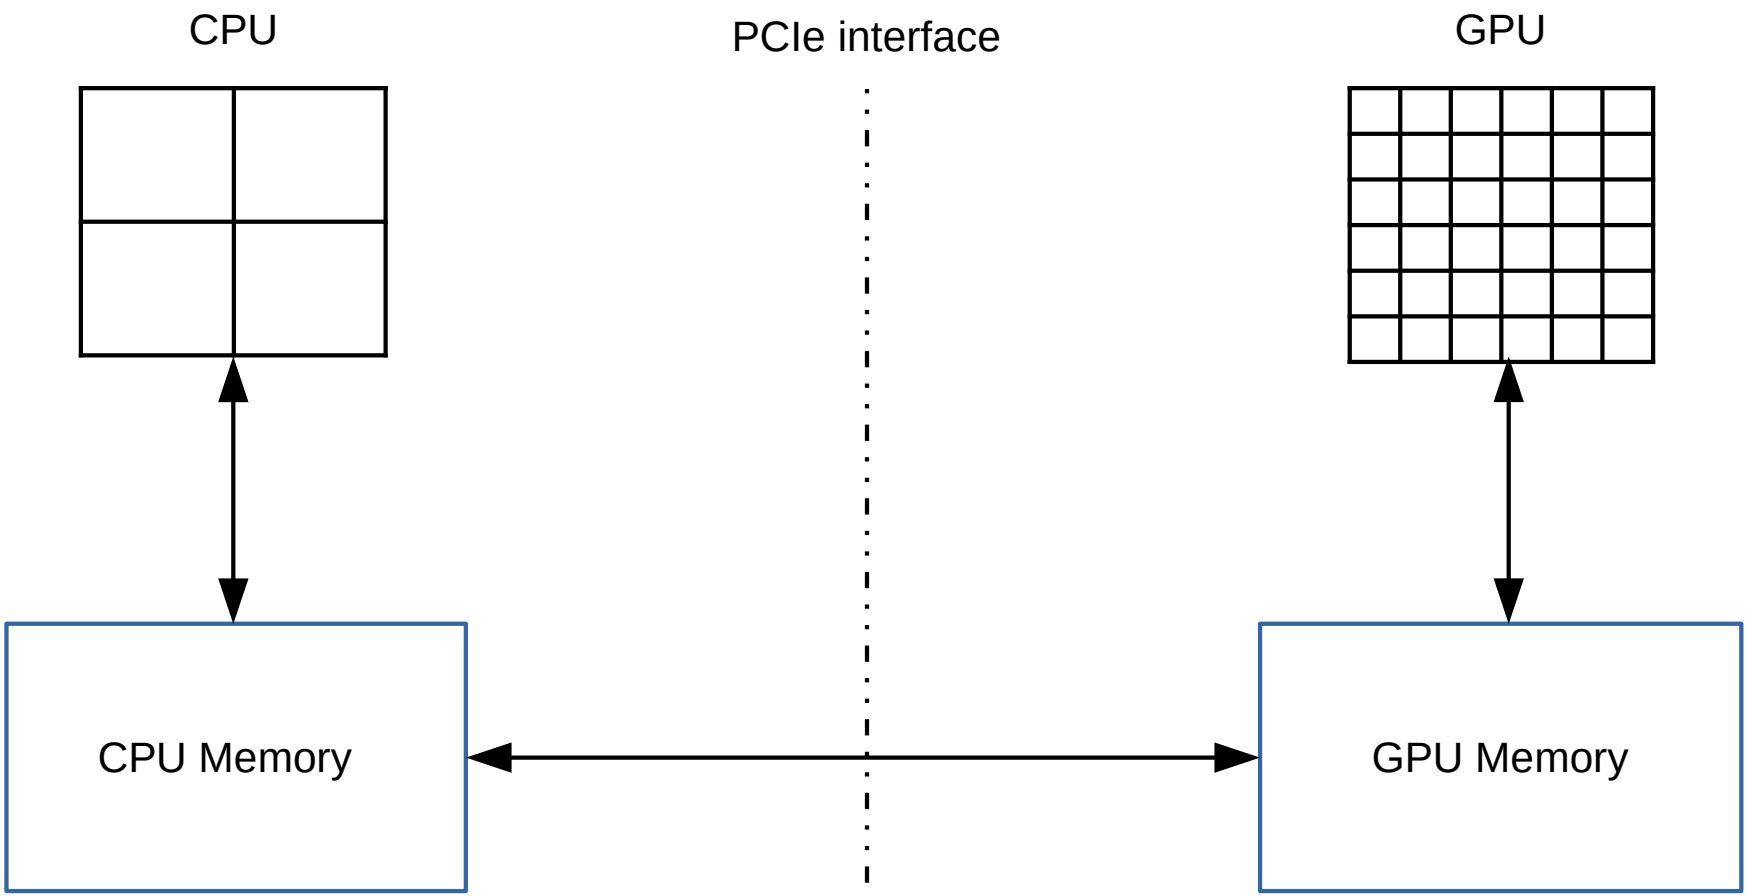

De-allocate the memories and terminate the program

Pseudo Code (new method)

1. Allocate and initialize Host Memory
2. Launch “kernel” on the Device – large number of threads execute in parallel on Device (Magic Step!)
3. De-allocate all the memory and terminate the program

Step 1

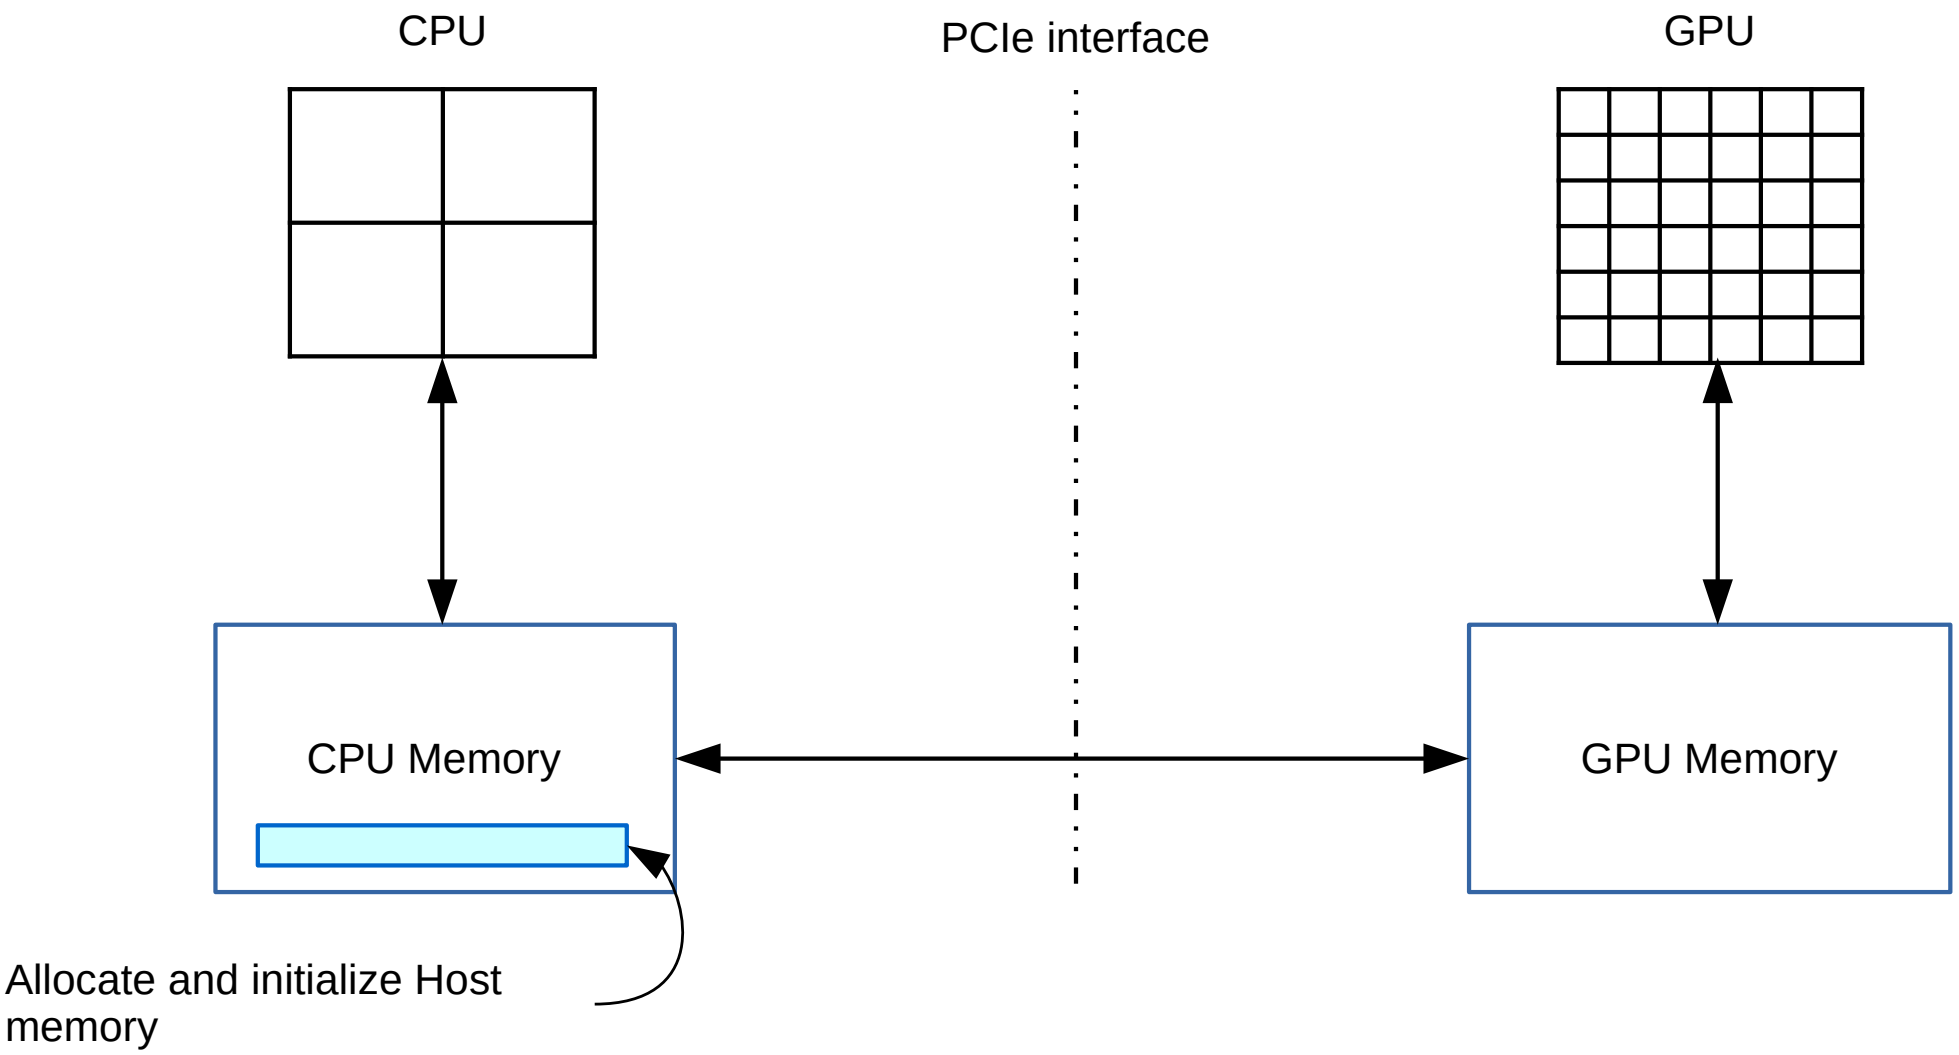
CPU
PCIe interface
GPU
CPU Memory
GPU Memory
Allocate and initialize Host memory

Step 2 (Magic Step!)

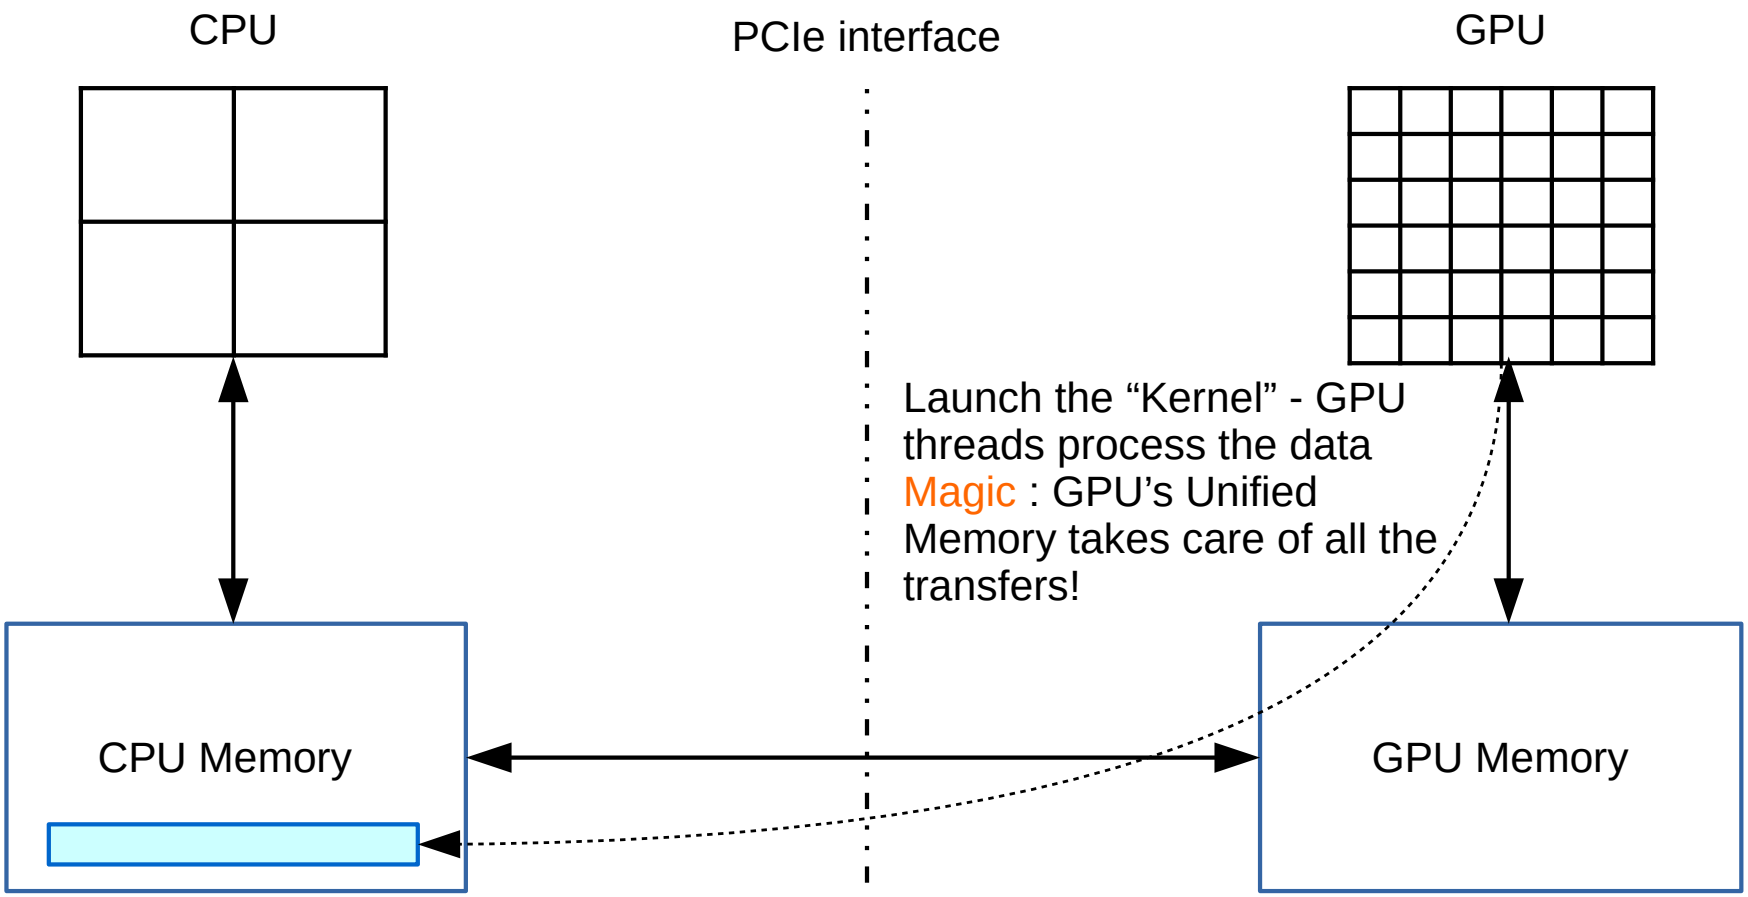

Step 3

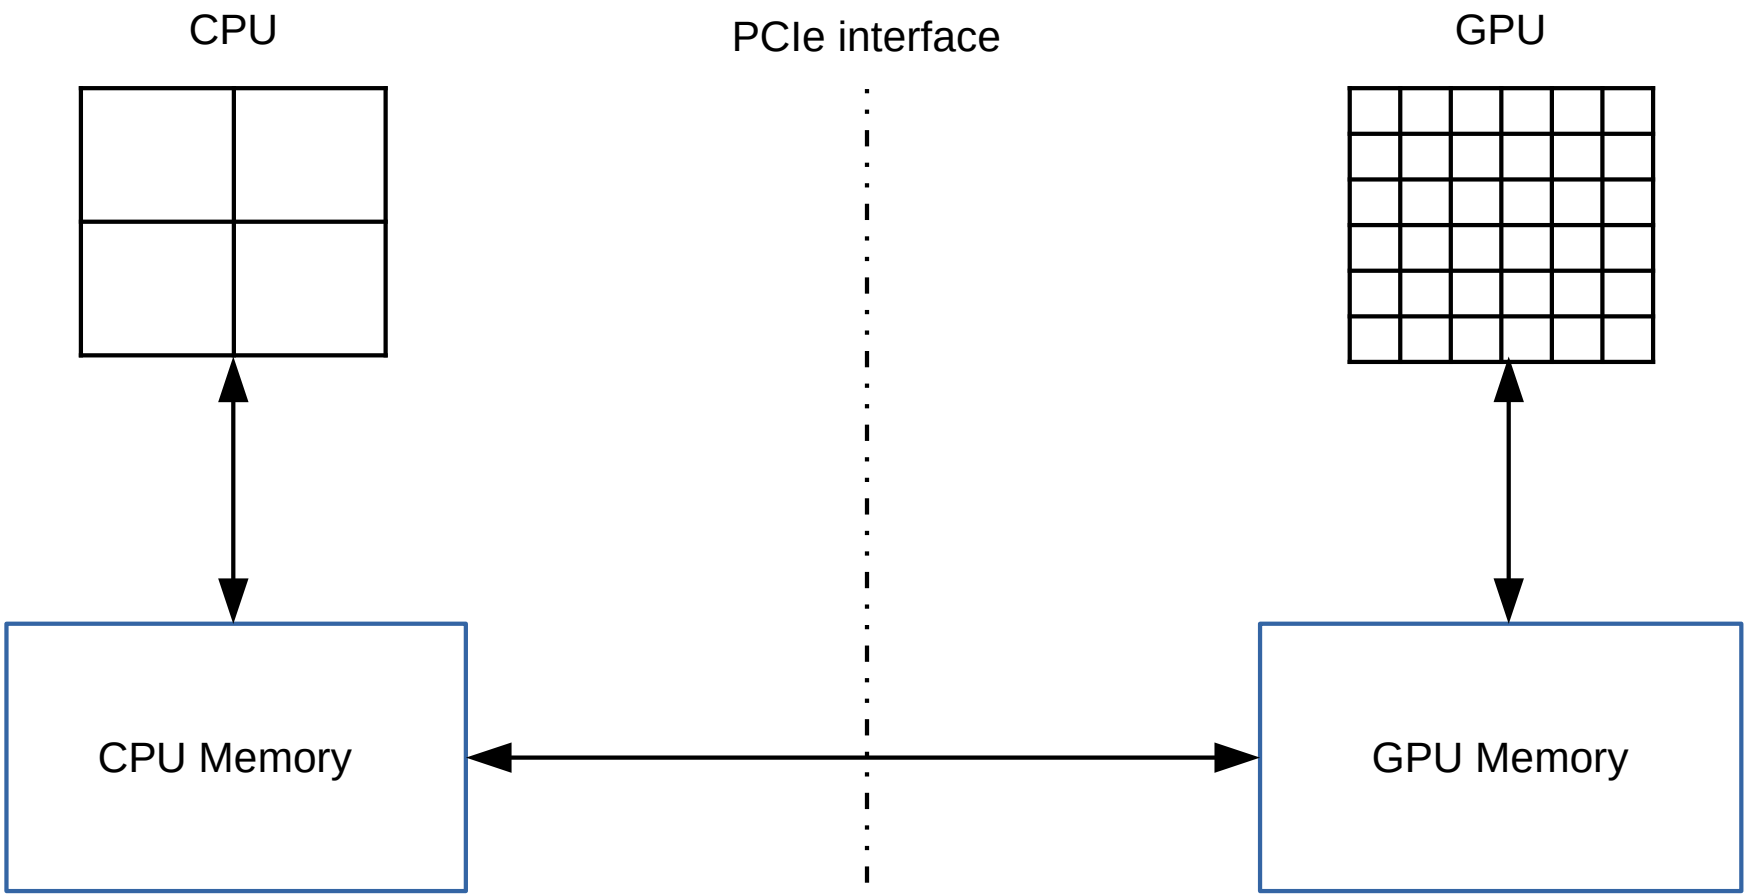

De-allocate the memory and terminate the program

Introduction to CUDA

CUDA

- Parallel computing platform and application programming interface (API)
- Nvidia Corp.- General Purpose Graphics Processing Unit (GPGPU)
- Supports most of the today's Nvidia's gaming cards.
- Platforms
  - Geforce : Desktop
  - Quadro : Workstation
  - Tesla : Datacenter
  - Tegra, Jetson, Drive : Embeded

Structure of a CUDA Code

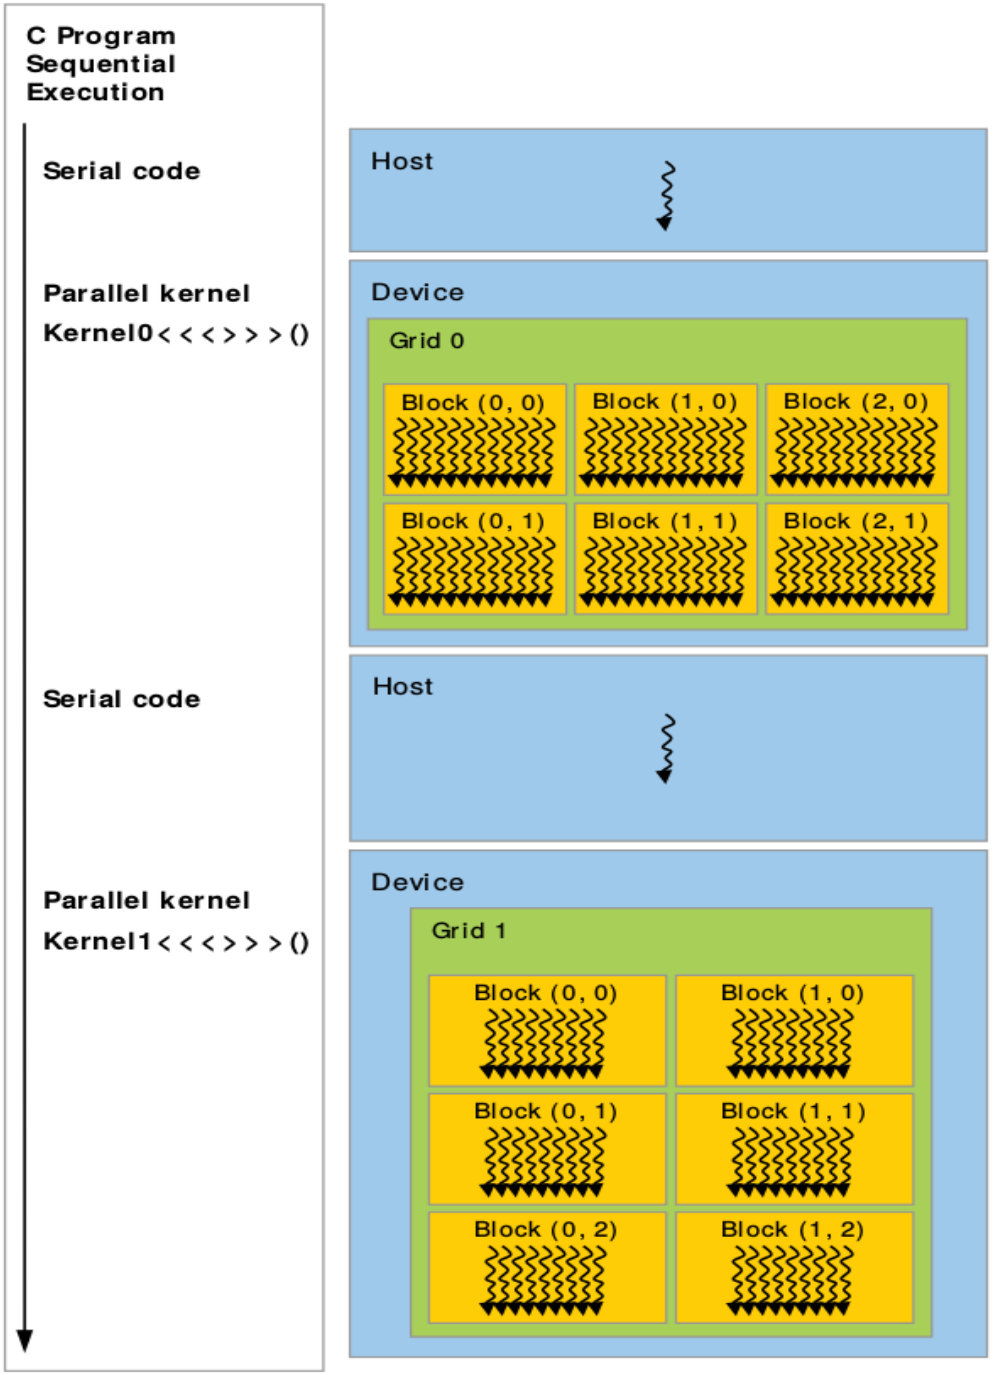

Source: CUDA C Programming Guide (NVIDIA)

Threads, Blocks and Grid(s)

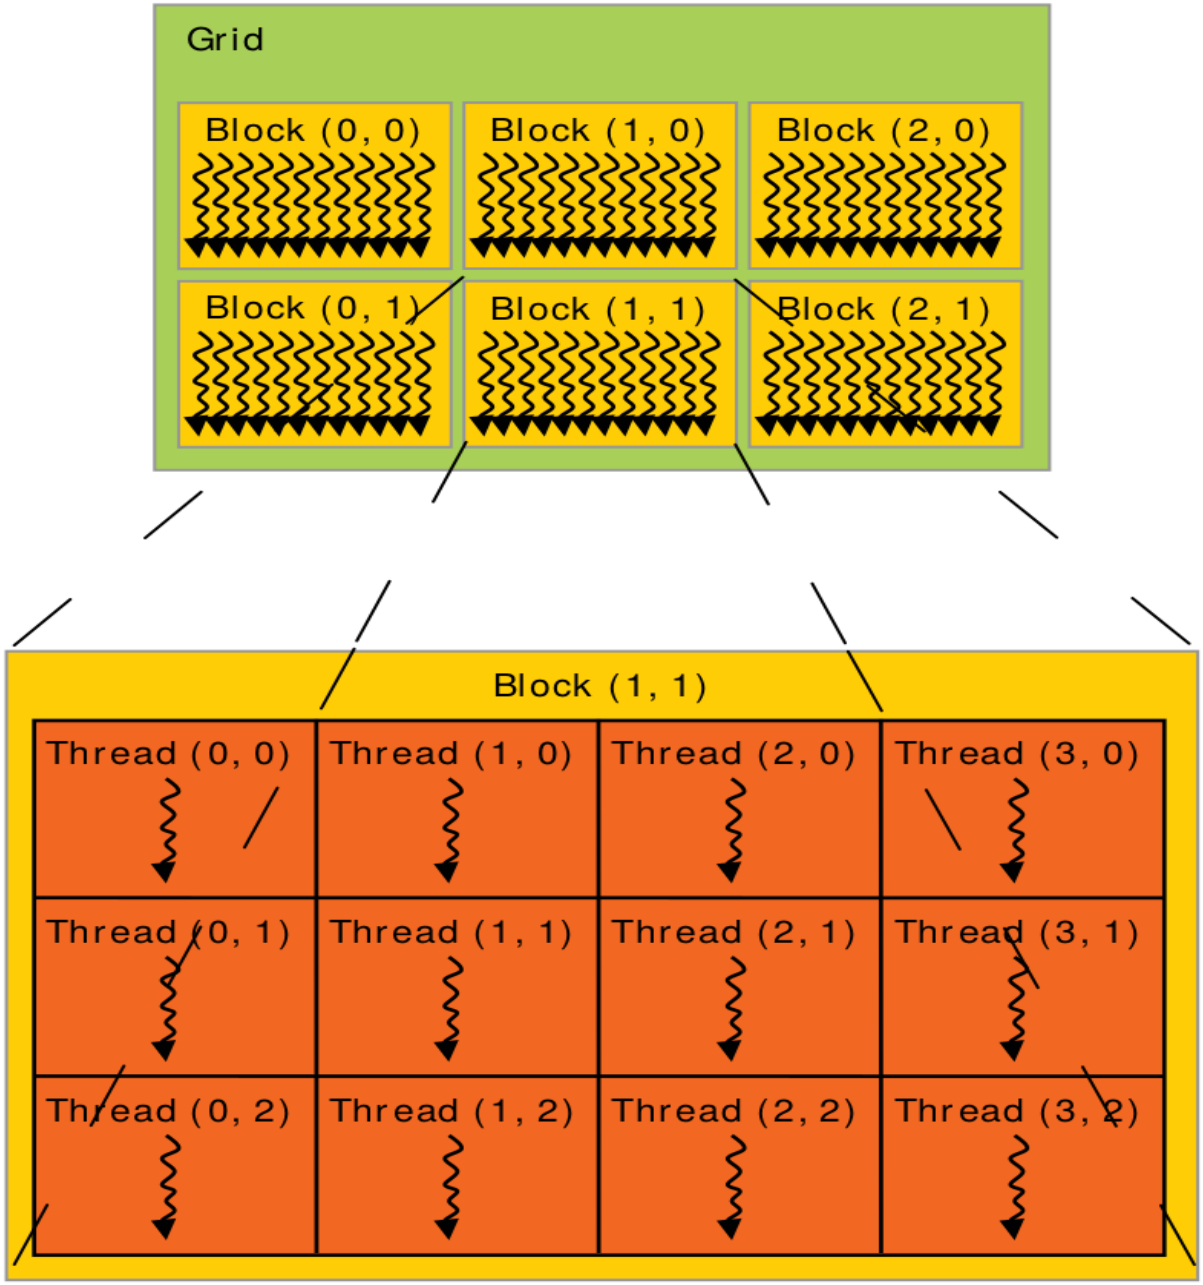

Source: CUDA C Programming Guide (NVIDIA)

Vector addition (old method)

```
int m[200], n[200], p[200],*md, *nd,*pd;
int size = 200 * sizeof(int);
// Initialize array m and array n
...
cudaMalloc(&md, size);
cudaMemcpy(md, m, size, cudaMemcpyHostToDevice);
cudaMalloc(&nd, size);
cudaMemcpy(nd, n, size, cudaMemcpyHostToDevice);
cudaMalloc(&pd, size);
arradd<<< 1,200 >>>(md,nd,pd);
cudaMemcpy(p, pd, size, cudaMemcpyDeviceToHost);
cudaFree(md);
cudaFree(nd);
cudaFree(pd);
```

“Kernel” function

```
__global__ void arradd(int* md, int* nd, int* pd)
{
int myid = threadIdx.x;

pd[myid] = md[myid] + nd[myid];
}
```

Questions?

Thank you.